

CSC411: Project 4

Due on Sunday, April 2, 2018

Ying Yang

March 29, 2018

Part 1

Environment

The grid is represented by horizontal, vertical and diagonal indices. The coordinates are the indices from 0 - 8. The win-set is represented by a fix matrix in which the first row represents 3 sets of horizontal rows for which a player wins the game, the second row represent the vertical column, and the third row represents the diagonal of the tictactoe.

The attributes turn represents whose turn to play the game (player1 or player 2). The attribute done means whether the game is over (when a player wins, or a tie occurs).

Play a game of TicTacToe by calling the step(), and render() methods.

Listing 1: Compute network function

```
env.step(0)
(array([1, 0, 0, 0, 0, 0, 0, 0, 0]), 'valid', False)
env.render()
x..
5   ...
   ...
   ====
env.step(2)
10  (array([1, 0, 2, 0, 0, 0, 0, 0, 0]), 'valid', False)
env.render()
x.o
   ...
   ...
   ====
15  env.step(4)
   (array([1, 0, 2, 0, 1, 0, 0, 0, 0]), 'valid', False)
env.render()
x.o
.x.
20  ...
   ...
   ====
env.step(8)
25  (array([1, 0, 2, 0, 1, 0, 0, 0, 2]), 'valid', False)
env.render()
x.o
.x.
..o
   ====
30  env.step(6)
   (array([1, 0, 2, 0, 1, 0, 1, 0, 2]), 'valid', False)
env.render()
x.o
.x.
x.o
35  ...
   ...
   ====
env.step(1)
   (array([1, 2, 2, 0, 1, 0, 1, 0, 2]), 'valid', False)
env.render()
xoo
40  .x.
```

45

```
x.o
====
env.step(3)
(array([1, 2, 2, 1, 1, 0, 1, 0, 2]), 'win', True)
env.render()
xoo
xx.
x.o
====
```

Part 2

Complete the implementation so that policy is a neural network with one hidden layer

Part 2 (a)

Listing 2: Policy implementation

```

class Policy(nn.Module):
    """
    The Tic-Tac-Toe Policy
    """
5   def __init__(self, input_size=27, hidden_size=64, output_size=9):
    super(Policy, self).__init__()
    # TODO
    self.linear_f1 = nn.Linear(input_size, hidden_size)
    self.linear_f2 = nn.Linear(hidden_size, output_size)
10
    def forward(self, x):
    # TODO
    h = F.relu(self.linear_f1(x))
    out = F.softmax(self.linear_f2(h))
15   return out

```

Part 2 (b)

Listing 3: Policy

```

policy = Policy()
state = np.array([1,0,1,2,1,0,1,0,1])
state = torch.from_numpy(state).long().unsqueeze(0)
state = torch.zeros(3,9).scatter_(0, state, 1).view(1, 27)
5   print(state)

```

Listing 4: output

```

Columns 0 to 12
0      0      1      0      0      0      1      0      1      0      1      0      1

Columns 13 to 25
5   0      1      0      1      0      1      0      0      0      1      0      0      0

Columns 26 to 26
0
[torch.FloatTensor of size 1x27]

```

State what each of the 27 dimensions mean

Part 2 (c)

Explain what the value in each dimension means. The value in each dimension means Is this policy stochastic or deterministic?

The policy is stochastic.

Part 3

Part 3a

Implement the compute_returns function

Listing 5: output

5

```
l = len(rewards)
rewards = np.array(rewards)
gammas = np.array([gamma ** (i) for i in range(l)])

G = []
for i in range(l):
    G.append(sum(rewards[i:] * gammas[:l - i]))
return G
```

Part 3b

Explain why can we not update weights in the middle of an episode

Part 4

Train the neural network using gradient descent (without momentum). Plot the learning curves. Display the weights going into each of the output units.

Optimization procedure:

We initialized the parameters as follows:

Listing 6: Parameter Initialization

```

# load weights and biases
snapshot = cPickle.load(open("snapshot50.pkl"))
W0 = snapshot["W0"] # dimension = (748, 300)
b0 = snapshot["b0"].reshape((300, 1))
5 W1 = snapshot["W1"] # dimension = (300, 10)
b1 = snapshot["b1"].reshape((10, 1))

# initiate weights, bias and alpha
initial_weight = dot(W0, W1) # (748, 10)
10 initial_bias = b1 # (10, 1)
alpha = 1e-5 # learning rate

EPS = 1e-10
prev_w = initial_weight - 10 * EPS
15 prev_b = initial_bias - 10 * EPS
W = initial_weight.copy()
b = initial_bias.copy()
iteration = 0
max_iteration = 4000

```

We initialized the weight to be the dot product of W0 and W1. We repeatedly experimented with different alpha values and observe the results. We first set alpha=0.001. The percentage of correctness turned out to be very low (with 25% of correctness). Then we decreased the learning rate to 0.00001. The correctness percentage then improved to 89%.

Gradient descent procedure:

Listing 7: Parameter Initialization

```

# apply gradient descent
while norm(W - prev_w) > EPS and norm(b - prev_b) > EPS and iteration < max_iteration:
    prev_w = W.copy()
    prev_b = b.copy()
5 W -= alpha * df_w(train_set.T, W, b, train_label.T)
b -= alpha * df_b(train_set.T, W, b, train_label.T)
    if iteration % 100 == 0:
        epoch.append(iteration)
        train_performance.append(performance(train_set.T, W, b, train_label.T))
10 test_performance.append(performance(test_set.T, W, b, test_label.T))
    iteration += 1

```

Learning curve is shown in the figure below:

Part 5

Write vectorized code that performs gradient descent with momentum, and use it to train your network. Plot the learning curves

Learning curve with momentum:

How new learning curves compare with gradient descent without momentum: Without the momentum, the performance of correct guesses increased slower than if momentum were used. For example, as we can observe in the learning curve, the performance of training set only increased to 92% after 800 iterations, whereas with momentum term set to 0.99, the performance correctness increased to the same percentage by only using 400 iterations, because it moved faster when gradient consistently points to one direction.

Code added in the gradient descent:

In the `train_neural_network` method, method signature was modified to take two extra terms: `type` and `momentum` term. The rest of code remains unchanged as in part 4.

Listing 8: Gradient descent with momentum

```
train_neural_network(image_prefix, type="None", momentum_term=0):  
  
    v_w = 0 # for momentum  
    v_b = 0  
    if type == "momentum":  
        v_w = momentum_term * v_w + alpha * df_w(train_set.T, W, b, train_label.T)  
        W -= v_w  
        v_b = momentum_term * v_b + alpha * df_b(train_set.T, W, b, train_label.T)  
        b -= v_b  
  
    if __name__ == "__main__":  
        momentum_term = 0.99  
        train_neural_network("part5", "momentum", momentum_term)
```

Part 6**Part 6a**

Plot the contour of the cost function

Part 6b**Part 6c****Part 6d**

Part 7

Part 8